

Machine Intelligence: Deep Learning

Uncertainty in DL models

Latest Development in TF

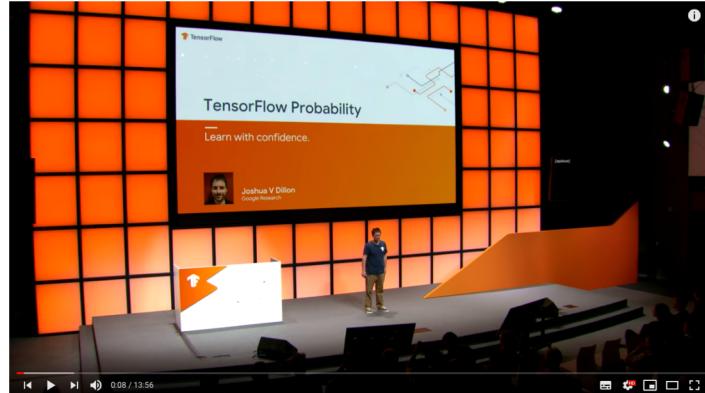
Probability

Oliver Dürr

Remark: Much of the material has been developed together with Elvis Murina and Beate Sick

Topics

- Sneak preview into TF Probability
- Demo:
 - Linear Regression (I) The usual probabilistic way
 - Linear Regression (II) Non-Fixed Variance
 - Linear Regression (III) Bayesian Way



[TF Probability Talk 6/3/2019](#)

Disclaimer: Fresh from the loom, might be a bit rough

Docker:

```
docker pull oduerr/tf_docker:cpu_r_tfp
```

Recap: Linear regression in statistics

Model for the conditional probability distribution

CPD: $Y_{X_i} = (Y|X_i) \sim N(\mu_{x_i}, \sigma^2)$

$$\begin{aligned} Y_x &\in \sim \\ \mu_x &\in \sim \end{aligned}$$

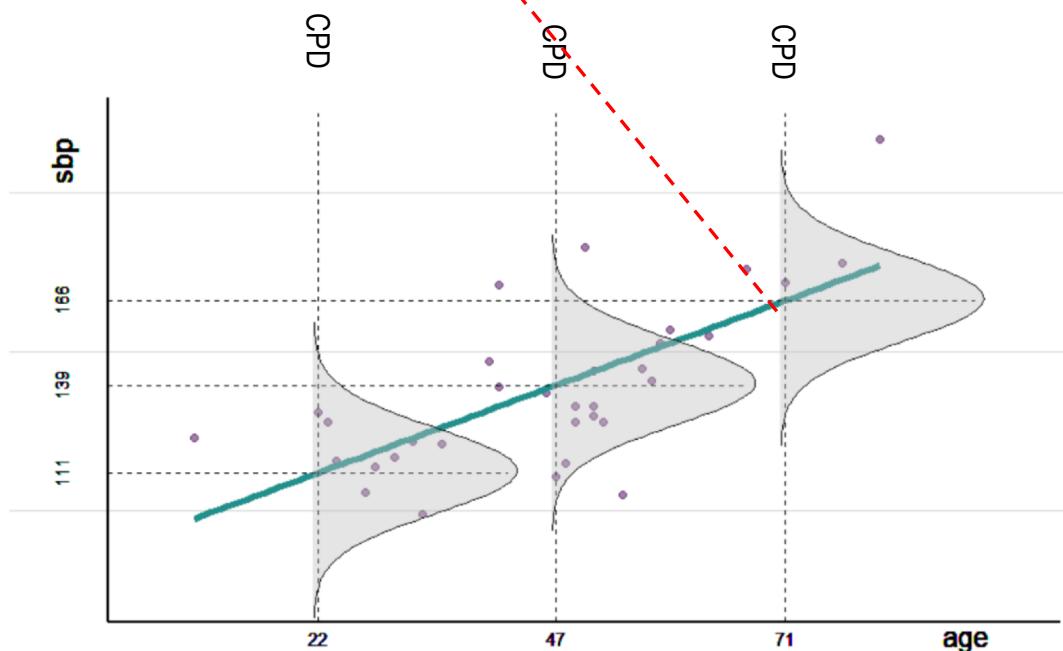
μ_x is modeled
 σ^2 is independent of the predictor values

$$y_i = \beta_0 + \beta_1 \cdot x_{i1} + \varepsilon_i$$

$$E(Y_{X_i}) = \mu_{x_i} = (\mu|X=x_i) = \beta_0 + \beta_1 \cdot x_{i1}$$

$$\text{Var}(Y_{X_i}) = \text{Var}(Y|X_i) = \text{Var}(\varepsilon_i) = \sigma^2$$

$$\varepsilon_i \text{ i.i.d. } \sim N(0, \sigma^2)$$



Demo also sigma

TFP: Neural Network Parametrizes Distribution

Model for the conditional probability distribution

CPD: $Y_{X_i} = (Y|X_i) \sim N(\mu_{x_i}, \sigma^2)$

Sigma is fixed

```
In [201]: def my_gauss(mu):
    return tfd.Normal(loc=mu, scale=1)

#Using the functional API for clarity
inputs = tf.keras.layers.Input(shape=(1,))
mu = tf.keras.layers.Dense(1)(inputs)
p_y = tfp.layers.DistributionLambda(my_gauss)(mu)

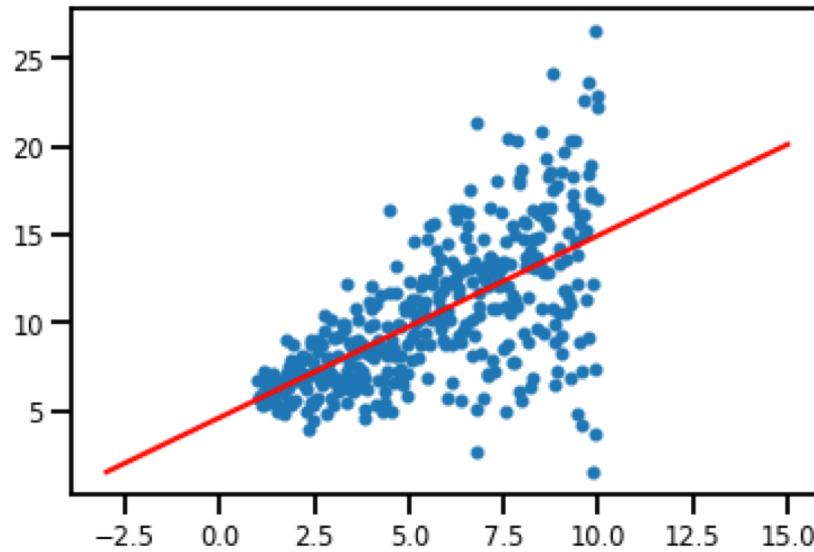
model = tf.keras.models.Model(inputs=inputs, outputs=p_y)
```

Also modeling sigma

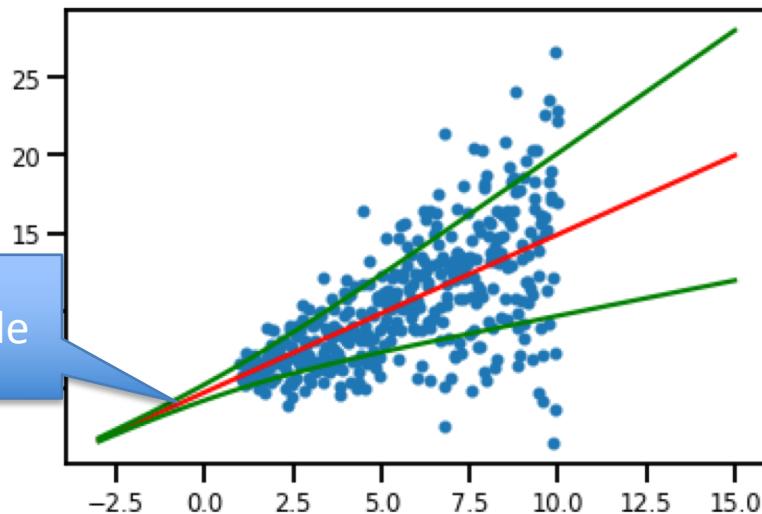
```
def my_gauss2(t):
    mu = t[:, :, 0] #First column mu
    scale = 1e-3 + tf.math.softplus(0.05 * t[:, :, 1:]) #Second column sigma
    return tfd.Normal(loc=mu, scale=scale)

#Using the functional API for clarity
inputs = tf.keras.layers.Input(shape=(1,))
t = tf.keras.layers.Dense(2)(inputs) #-- 2 outputs
p_y = tfp.layers.DistributionLambda(my_gauss2)(t)
model = tf.keras.models.Model(inputs=inputs, outputs=p_y)
```

Results



No modelling of sigma

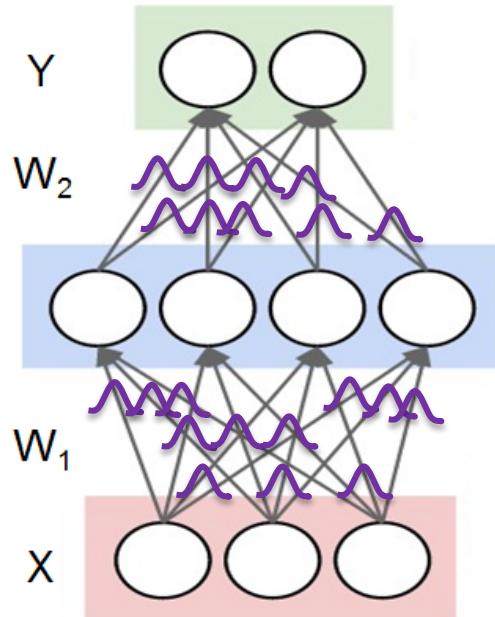


Modelling of sigma

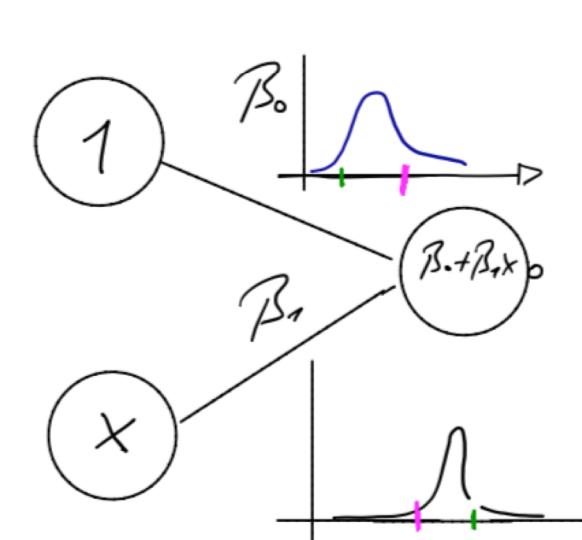
Questionable

Bayesian DL

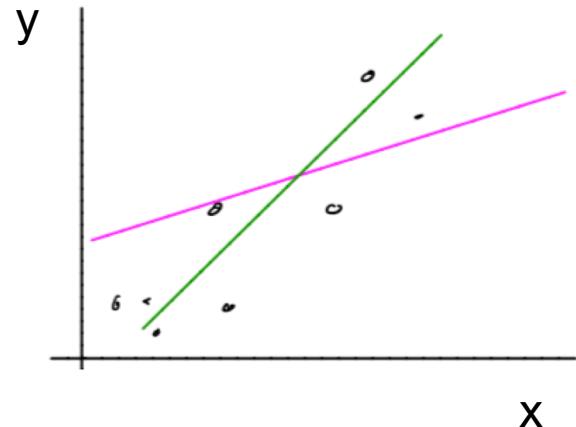
Modelling Uncertainty: Linear Bayes



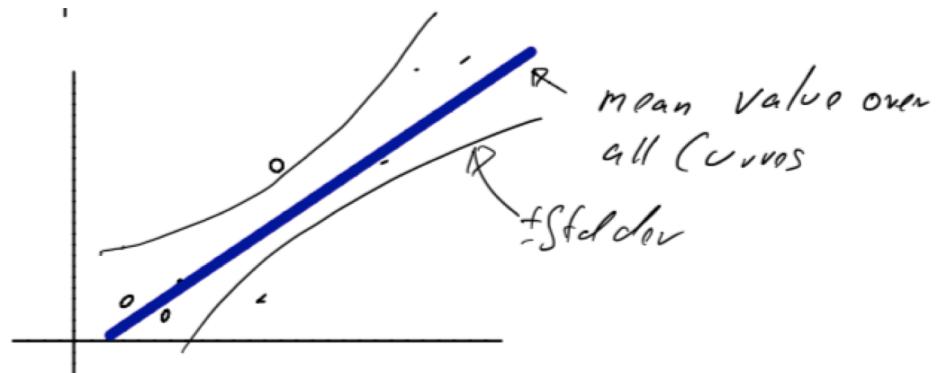
Bayesian NN



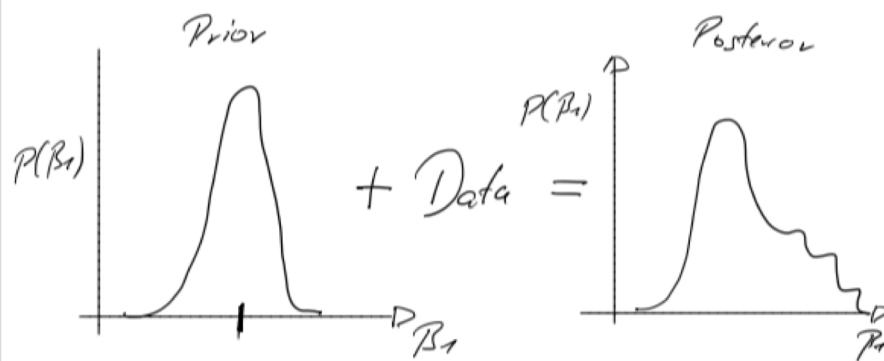
2 Examples



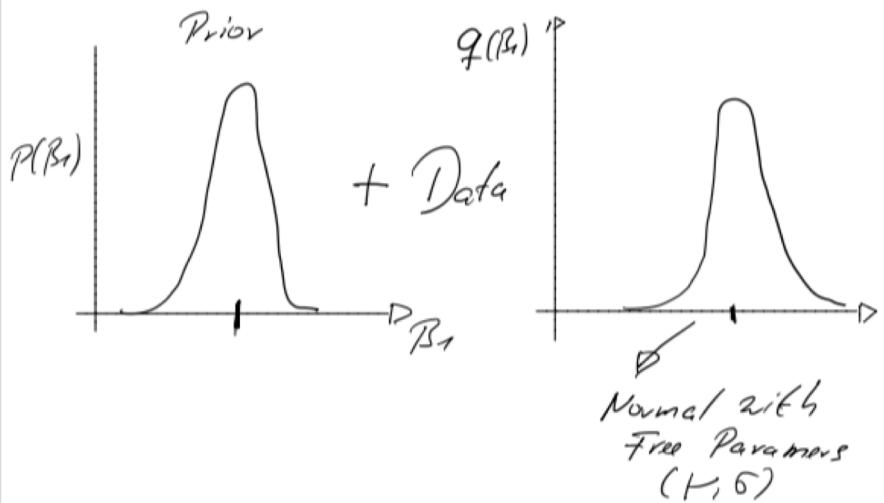
Averages over all



Training Using the variational approximation



“Normal” Bayes



“Variational Approximation” Bayes

R code for Variational Bayes

Before: no

```
In [201]: def my_gauss(mu):
    return tfd.Normal(loc=mu, scale=1)

#Using the functional API for clarity
inputs = tf.keras.layers.Input(shape=(1,))
mu = tf.keras.layers.Dense(1)(inputs)
p_y = tfp.layers.DistributionLambda(my_gauss)(mu)

model = tf.keras.models.Model(inputs=inputs, outputs=p_y)
```

```
#Using the functional API for clarity
inputs = tf.keras.layers.Input(shape=(1,))
tt = tfp.layers.DenseVariational(1, posterior_mean_field, prior_trainable)(inputs)
p_y = tfp.layers.DistributionLambda(my_gauss)(tt)

model = tf.keras.models.Model(inputs=inputs, outputs=p_y)
```

